

## Question 4 Image Classification (60 points) – Competition

In this exercise, your task is to create a Convolutional Neural Network (CNN) based image classifier for the CIFAR-10 dataset of images.

To solve the image classification problem using CIFAR-10 dataset, I went through online resources and followed the architecture from Stanford CS231n: Convolutional Neural Networks for Visual Recognition, lecture notes.

### Network Architecture:

Three main types of layers are used to build ConvNet architectures: **Convolutional Layer**, **Pooling Layer**, and **Fully-Connected Layer**. These layers are stacked to form a full ConvNet architecture. The sequence is as follows:

[INPUT - CONV - RELU - POOL - FC]

For example : if the input size is  $32 \times 32 \times 3$ , the following steps are performed to apply ConvNet.

- INPUT holds raw pixel values of the image, in this case an image of width 32, height 32, and with three color channels R, G, B.
- CONV layer computes the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume. This results in volume such as  $[32 \times 32 \times 3]$  as we decided to use 3 filters.
- RELU layer applies an elementwise activation function, such as the  $\max(0, x)$  thresholding at zero. This leaves the size of the volume unchanged ( $[32 \times 32 \times 3]$ ).
- POOL layer performs a down sampling operation along the spatial dimensions (width, height), resulting in volume such as  $[16 \times 16 \times 12]$ .
- FC (i.e. fully-connected) layer computes the class scores, resulting in volume of size  $[1 \times 1 \times 10]$ , where each of the 10 numbers correspond to a class score, such as among the 10 categories of CIFAR-10. As with ordinary Neural Networks and as the name implies, each neuron in this layer will be connected to all the numbers in the previous volume.

**85% testing accuracy** is achieved with 150 epoch. Some of the other details of this architecture are summarized in the table.

Loss function	Optimization	Parameters
Cross-entropy	Stochastic Gradient Descent	Learning rate=0.001, weight_decay= 1E-5, momentum=0.9

To deal with overfitting problem, I have used dropout, L2 regularization, and batch normalization.

**Dropout** randomly drop units (along with their connections) from the neural network during training. The reduction in number of parameters in each step of training has effect of regularization.

**L2 regularization** allows to apply penalties on layer parameters during optimization. During gradient descent parameter update, the L2 regularization is applied using weight\_decay parameter in SGD.

**BatchNormalization** normalizes the activation of the previous layer at each batch, i.e. applies a transformation that maintains the mean activation close to 0 and the activation standard deviation close to 1.

Architecture:

